

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: df = pd.read_csv(r"C:\Users\LIKHITH\OneDrive\Desktop\Data science\Data_files\Salary
```

```
In [4]: df
```

Out[4]:

	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891
5	2.9	56642
6	3.0	60150
7	3.2	54445
8	3.2	64445
9	3.7	57189
10	3.9	63218
11	4.0	55794
12	4.0	56957
13	4.1	57081
14	4.5	61111
15	4.9	67938
16	5.1	66029
17	5.3	83088
18	5.9	81363
19	6.0	93940
20	6.8	91738
21	7.1	98273
22	7.9	101302
23	8.2	113812
24	8.7	109431
25	9.0	105582
26	9.5	116969
27	9.6	112635
28	10.3	122391
29	10.5	121872

```
In [5]: df.head()
```

```
Out[5]:
```

	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891

```
In [7]: df.shape
```

```
Out[7]: (30, 2)
```

```
In [8]: df.columns
```

```
Out[8]: Index(['YearsExperience', 'Salary'], dtype='object')
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: YearsExperience    0  
Salary                  0  
dtype: int64
```

```
In [10]: df.dtypes
```

```
Out[10]: YearsExperience    float64  
Salary                  int64  
dtype: object
```

```
In [11]: X=df.drop('YearsExperience',axis=1)  
y=df['Salary']
```

```
In [12]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test=train_test_split(X, y, random_state=1234, test_siz
```

```
In [13]: X_train.shape, X_test.shape
```

```
Out[13]: ((21, 1), (9, 1))
```

```
In [14]: y_train.shape, y_test.shape
```

```
Out[14]: ((21,), (9,))
```

```
In [15]: df.shape
```

```
Out[15]: (30, 2)
```

```
In [16]: X_train
```

Out[16]:

Salary**13** 57081**22** 101302**24** 109431**0** 39343**2** 37731**27** 112635**26** 116969**18** 81363**5** 56642**16** 66029**25** 105582**11** 55794**9** 57189**17** 83088**29** 121872**20** 91738**12** 56957**21** 98273**6** 60150**19** 93940**15** 67938In [17]: `y_train`

```
Out[17]: 13      57081
          22     101302
          24     109431
           0      39343
           2      37731
          27     112635
          26     116969
          18      81363
           5      56642
          16     66029
          25     105582
          11     55794
           9      57189
          17     83088
          29     121872
          20     91738
          12     56957
          21     98273
           6      60150
          19     93940
          15     67938
          Name: Salary, dtype: int64
```

```
In [18]: X_test
```

```
Out[18]:
```

	Salary
7	54445
10	63218
4	39891
1	46205
28	122391
8	64445
3	43525
23	113812
14	61111

```
In [19]: y_test
```

```
Out[19]: 7      54445
          10     63218
          4      39891
          1      46205
          28    122391
          8      64445
          3      43525
          23    113812
          14     61111
          Name: Salary, dtype: int64
```

```
In [20]: X_train.ndim
```

```
Out[20]: 2
```

```
In [21]: from sklearn.linear_model import LinearRegression
          LR=LinearRegression()
          LR.fit(X_train,y_train)
```

```
Out[21]: ▼ LinearRegression ⓘ ?
          LinearRegression()
```

```
In [22]: # Model predictions happens X_test
          y_predictions=LR.predict(X_test)
```

```
In [23]: y_predictions
```

```
Out[23]: array([ 54445.,  63218.,  39891.,  46205., 122391.,  64445.,  43525.,
                113812.,  61111.])
```

```
In [24]: y_test.shape,y_predictions.shape
```

```
Out[24]: ((9,), (9,))
```

```
In [25]: X_test
```

Out[25]:

	Salary
7	54445
10	63218
4	39891
1	46205
28	122391
8	64445
3	43525
23	113812
14	61111

```
In [26]: X_test.iloc[0] # series
# In order to pass a test sample to a model
# we need to pass a list of values
# or array of values
# tuple of values
X_test.iloc[0].values
```

Out[26]: array([54445])

```
In [27]: LR.predict([X_test.iloc[0].values,
X_test.iloc[1].values])
```

Out[27]: array([54445., 63218.])

```
In [28]: ip1=[5]
LR.predict([ip1])
```

Out[28]: array([5.])

```
In [29]: X_test.shape,y_test.shape,y_predictions.shape
```

Out[29]: ((9, 1), (9,), (9,))

```
In [30]: test_data=X_test
test_data['y_actual']=y_test
test_data['y_predictions']=y_predictions
test_data
```

Out[30]:

	Salary	y_actual	y_predictions
7	54445	54445	54445.0
10	63218	63218	63218.0
4	39891	39891	39891.0
1	46205	46205	46205.0
28	122391	122391	122391.0
8	64445	64445	64445.0
3	43525	43525	43525.0
23	113812	113812	113812.0
14	61111	61111	61111.0

```
In [31]: # y_test is series
# y_predictions is numpy array values
print(y_test.values[:5]) # float 5. means 5.0
print(y_predictions[:5])
```

```
[ 54445  63218  39891  46205 122391]
[ 54445.  63218.  39891.  46205. 122391.]
```

```
In [32]: # RMSE
# MSE
# MAE
# R-square
from sklearn.metrics import r2_score, mean_squared_error
```

```
In [33]: R2=r2_score(y_test,y_predictions)
MSE=mean_squared_error(y_test,y_predictions)
#MSE**(1/2)
RMSE=np.sqrt(MSE)
#accuracy_score(y_test,y_predictions) # it is a regression tech
print("R-sqaure:",R2)
print("MSE:",MSE)
print("RMSE:",RMSE)
```

```
R-sqaure: 1.0
MSE: 6.470390569303684e-23
RMSE: 8.043873798925294e-12
```

```
In [35]: s=0
for i in range(len(y_test)):
    v1=y_test.values[i]-y_predictions[i]
    v2=v1**2
    s=s+v2
print(s/len(y_test))
```

```
6.470390569303684e-23
```

```
In [36]: LR.coef_
```



```
print("The coeffiecnt of Years_of_experience is:", LR.coef_)
```

The coeffiecnt of Years_of_experience is: [1.]

```
In [37]: LR.intercept_
```

```
Out[37]: np.float64(-1.4551915228366852e-11)
```

```
In [38]: X_train.columns
```

```
Out[38]: Index(['Salary'], dtype='object')
```

```
In [39]: #Regression_equation=LR.intercept_+LR.coef_ * col namee
#Regression_equation
y=-1.45+1-.*Salary
```

Cell In[39], line 3

y=-1.45+1-.*Salary

SyntaxError: invalid syntax

```
In [40]: from sklearn.feature_selection import VarianceThreshold
vt=VarianceThreshold(threshold=0)
# Threshold variance value
# we want to drop the feaure based on threshold
vt.fit(df)
```

```
Out[40]: ▼ VarianceThreshold ⓘ ?
VarianceThreshold(threshold=0)
```

```
In [41]: dir(vt)
```

```
Out[41]: ['__abstractmethods__',
          '__annotations__',
          '__class__',
          '__delattr__',
          '__dict__',
          '__dir__',
          '__doc__',
          '__eq__',
          '__firstlineno__',
          '__format__',
          '__ge__',
          '__getattr__',
          '__getstate__',
          '__gt__',
          '__hash__',
          '__init__',
          '__init_subclass__',
          '__le__',
          '__lt__',
          '__module__',
          '__ne__',
          '__new__',
          '__reduce__',
          '__reduce_ex__',
          '__repr__',
          '__setattr__',
          '__setstate__',
          '__sizeof__',
          '__sklearn_clone__',
          '__sklearn_tags__',
          '__static_attributes__',
          '__str__',
          '__subclasshook__',
          '__weakref__',
          '_abc_impl',
          '_build_request_for_signature',
          '_check_feature_names',
          '_check_n_features',
          '_doc_link_module',
          '_doc_link_template',
          '_doc_link_url_param_generator',
          '_get_default_requests',
          '_get_doc_link',
          '_get_metadata_request',
          '_get_param_names',
          '_get_support_mask',
          '_get_tags',
          '_more_tags',
          '_parameter_constraints',
          '_repr_html_',
          '_repr_html_inner',
          '_repr_mimebundle_',
          '_sklearn_auto_wrap_output_keys',
          '_transform',
          '_validate_data',
          '_validate_params',
```

```
'feature_names_in_',  
'fit',  
'fit_transform',  
'get_feature_names_out',  
'get_metadata_routing',  
'get_params',  
'get_support',  
'inverse_transform',  
'n_features_in_',  
'set_output',  
'set_params',  
'threshold',  
'transform',  
'variances_']
```

```
In [43]: vt.variances_
```

```
Out[43]: array([7.78515556e+00, 8.46600000e+04])
```

```
In [44]: vt.get_support()
```

```
Out[44]: array([ True,  True])
```

```
In [45]: vt.get_params()
```

```
Out[45]: {'threshold': 0}
```

```
In [46]: vt.threshold
```

```
Out[46]: 0
```

```
In [47]: cols=vt.get_feature_names_out()  
# the above syntax gives the column names  
# These fetaure only we want include  
df[cols]
```

Out[47]:

	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891
5	2.9	56642
6	3.0	60150
7	3.2	54445
8	3.2	64445
9	3.7	57189
10	3.9	63218
11	4.0	55794
12	4.0	56957
13	4.1	57081
14	4.5	61111
15	4.9	67938
16	5.1	66029
17	5.3	83088
18	5.9	81363
19	6.0	93940
20	6.8	91738
21	7.1	98273
22	7.9	101302
23	8.2	113812
24	8.7	109431
25	9.0	105582
26	9.5	116969
27	9.6	112635
28	10.3	122391
29	10.5	121872

```
In [51]: path=r"C:\Users\LIKHITH\OneDrive\Desktop\Data science\Data_files\Salary_Data.csv"
df=pd.read_csv(path)
df.head()
from sklearn.feature_selection import VarianceThreshold
vt=VarianceThreshold(threshold=0)
### Make sure before fitting the dataframe , do not include output column
X=df.drop('YearsExperience',axis=1)
# X it self a data frame
vt.fit(X)
vt.variances_
vt.get_support()
cols=vt.get_feature_names_out()
X[cols]
```

Out[51]:

Salary

0	39343
1	46205
2	37731
3	43525
4	39891
5	56642
6	60150
7	54445
8	64445
9	57189
10	63218
11	55794
12	56957
13	57081
14	61111
15	67938
16	66029
17	83088
18	81363
19	93940
20	91738
21	98273
22	101302
23	113812
24	109431
25	105582
26	116969
27	112635
28	122391
29	121872

```
In [52]: from statsmodels.api import OLS
OLS(y_train,X_train).fit().summary()
```

Out[52]:

OLS Regression Results

Dep. Variable:	Salary	R-squared (uncentered):	1.000
Model:	OLS	Adj. R-squared (uncentered):	1.000
Method:	Least Squares	F-statistic:	3.694e+32
Date:	Tue, 06 Jan 2026	Prob (F-statistic):	3.81e-314
Time:	17:57:26	Log-Likelihood:	488.13
No. Observations:	21	AIC:	-974.3
Df Residuals:	20	BIC:	-973.2
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Salary	1.0000	5.2e-17	1.92e+16	0.000	1.000	1.000

Omnibus: 3.403 **Durbin-Watson:** 0.280

Prob(Omnibus): 0.182 **Jarque-Bera (JB):** 2.287

Skew: 0.627 **Prob(JB):** 0.319

Kurtosis: 1.979 **Cond. No.** 1.00

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [53]: import pickle
pickle.dump(LR,
open('YearsExperience_model.pkl','wb'))
```

```
In [54]: # Loading model to compare the result
model=pickle.load(open('YearsExperience_model.pkl','rb'))
model
```

Out[54]:

▼ LinearRegression ⓘ ?

LinearRegression()

In []: